

SOLID.AI Framework

A Formal Specification for Strategic, Organized, Layered, Intelligent,
Data-Driven Artificial Intelligence

Whitepaper v1.0 — Stable

Published: December 2025

License: MIT

A comprehensive architectural specification for building AI-native organizations that scale human intelligence through structured collaboration between people and artificial intelligence. This whitepaper provides the complete technical specification, architectural patterns, and governance principles required to implement production-ready AI-native systems.

Abstract

Status: Version: 1.0

Citation

If you use SOLID.AI in your research or project, please cite:

```
@dataset{solidai_zenodo_2025,
  title      = {SOLID.AI Framework – Whitepaper v1.0},
  author     = {Freitas, Gustavo},
  year       = 2025,
  month      = november,
  publisher  = {Zenodo},
  doi        = 10.5281/zenodo.17765515,
  url        = https://zenodo.org/records/17765515
}
```

APA:

Freitas, G. (2025). SOLID.AI Framework – Whitepaper v1.0 [Dataset]. Zenodo. <https://doi.org/10.5281/zenodo.17765515>

IEEE:

G. Freitas, "SOLID.AI Framework – Whitepaper v1.0," Zenodo, Nov. 2025.
doi: 10.5281/zenodo.17765515

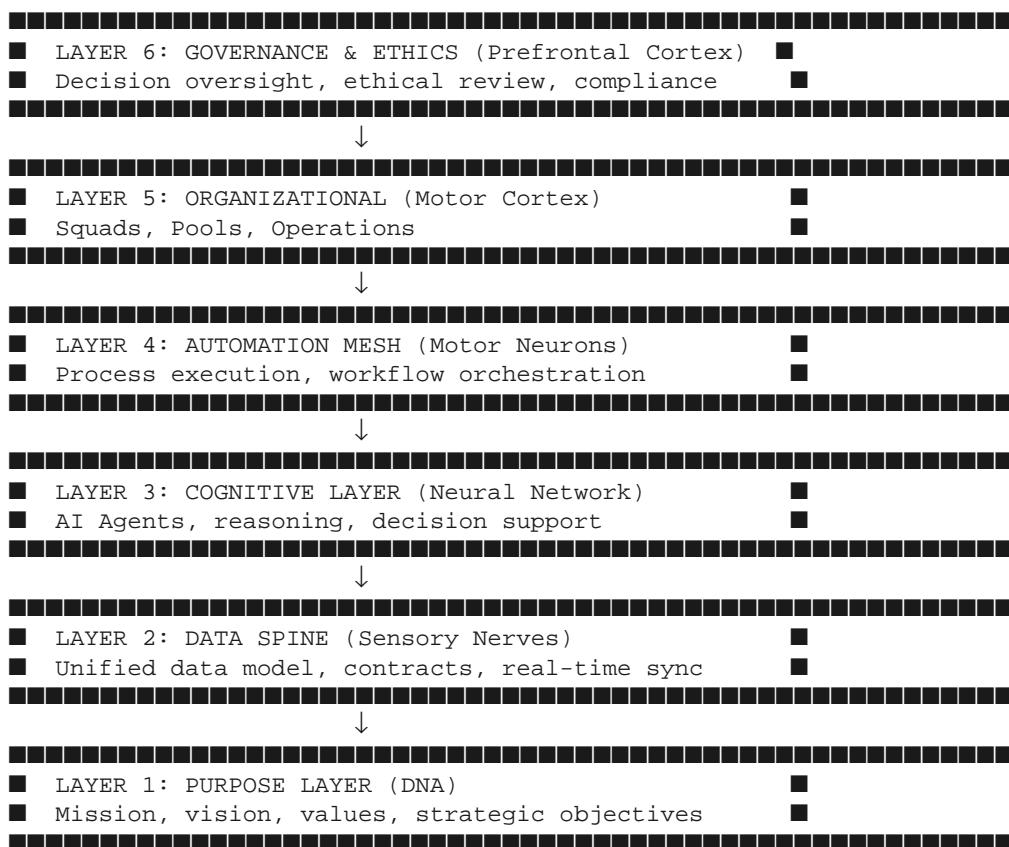
Navigation: [Architecture →](#)

Architecture

Status: Version: 1.0

Six-Layer Architecture

As shown in Figure 1 (see [Diagrams](#)), SOLID.AI employs a biological-inspired architecture analogous to an organizational nervous system:



See: [Interactive Layer Model Diagram](#) for detailed visualization

Layer 1: Purpose Layer (DNA)

Biological Analogy: DNA encoding the organism's fundamental blueprint

Function: Defines the organization's immutable core identity and strategic direction

Components:

- Mission statement
- Vision and strategic goals
- Core values and principles
- Success metrics (OKRs)
- Ethical boundaries

Key Characteristics:

- Rarely changes (only through formal RFC process)
- Informs all decisions across layers
- Accessible to all humans and AI agents
- Machine-readable format (YAML/JSON)

Layer 2: Data Spine (Sensory Nerves)

Biological Analogy: Sensory nervous system transmitting information to the brain

Function: Unified, real-time data infrastructure serving as single source of truth

As shown in Figure 3 (see [Diagrams](#)), the Data Spine is designed to meet stringent Service Level Objectives: **P95 latency < 5s, availability ≥ 99.9%, data freshness < 60s** (target specification).

Components:

- Canonical data models
- Data contracts between systems
- Event-driven synchronization
- Data quality monitoring
- Analytics and metrics dashboards

Key Characteristics:

- Schema-first design with strict contracts
- Real-time propagation (P95 latency < 5s)
- Immutable event logs (audit trail)
- Bi-directional sync across all systems

Architectural Foundation: The Data Spine implements data mesh principles defined by Dehghani[^4]: data as a product, domain ownership, self-serve data platform, and federated computational governance. Systematic research[^5] validates distributed data backbones with federated governance as essential for modern organizational data infrastructure, functioning as SOLID.AI's "organizational nervous system."

[^4]: Dehghani, Z. (2022). *Data Mesh Principles and Logical Architecture.* <https://martinfowler.com/articles/data-mesh-principles.html>

[^5]: Goedegebuure, A., Burnay, C., & van der Werf, J. M. (2023). *Data Mesh: A Systematic Gray Literature Review.* arXiv:2304.01062. <https://arxiv.org/abs/2304.01062>

See: [Specification → Data Spine](#) | [Data Spine Topology Diagram](#)

Layer 3: Cognitive Layer (Neural Network)

Biological Analogy: Brain processing information and generating insights

Function: AI agents providing reasoning, decision support, and autonomous actions

Components:

- AI Agent definitions (capabilities, constraints, interfaces)
- Reasoning engines (LLM orchestration)
- Context management (memory, session state)
- Decision logs (transparency)

Agent Types:

- **Analytical Agents:** Data analysis, pattern recognition, forecasting

- **Operational Agents:** Process execution, workflow orchestration
- **Advisory Agents:** Strategic recommendations, risk assessment
- **Collaborative Agents:** Team coordination, meeting facilitation

Research Validation: MIT Sloan research demonstrates AI tends to complement rather than replace human work, with deployment strategy (augmentation vs. replacement) being a strategic leadership decision[^6]. Harvard Business Review identifies hybrid human-AI teams as generating greatest value when processes and roles are redesigned for collaboration, not replacement[^7]—the foundation of SOLID.AI's Human-AI Collaboration Loop (Figure 4).

[^6]: MIT Sloan (2025). *New MIT Sloan Research Suggests AI is More Likely to Complement, Not Replace, Human Workers.* <https://mitsloan.mit.edu/press/new-mit-sloan-research-suggests-ai-more-likely-to-complement-not-replace-human-workers>

[^7]: Wilson, H. J., & Daugherty, P. R. (2018). *Collaborative Intelligence: Humans and AI Are Joining Forces.* Harvard Business Review. <https://hbr.org/2018/07/collaborative-intelligence-humans-and-ai-are-joining-forces>

See: [Specification → Cognitive Layer](#)

Layer 4: Automation Mesh (Motor Neurons)

Biological Analogy: Motor nervous system executing coordinated movements

Function: Process execution layer translating decisions into actions

Figure 2 — Automation Mesh Execution Model

As shown in Figure 2 (see [Diagrams](#)), the Automation Mesh coordinates all AI-driven actions through event-driven orchestration connecting agents, business services, and external systems.

Components:

- SIPOC process definitions
- Workflow orchestration (temporal.io, Airflow)
- Integration adapters (APIs, webhooks)
- Monitoring and observability

Key Patterns:

- **SIPOC Automation:** Supplier → Input → Process → Output → Customer
- **Event-Driven Workflows:** Trigger → Validate → Execute → Verify
- **Human-in-the-Loop:** Approval gates for critical decisions

Orchestration Pattern: SOLID.AI combines centralized orchestration with event-based choreography^[^8], leveraging event-driven architecture for service decoupling, resilience, and scalability^[^9]—enabling the Automation Mesh to coordinate AI agents, business services, and human workflows without brittle point-to-point integrations.

[^8]: Camunda (2023). *Orchestration vs. Choreography in Microservices*. <https://camunda.com/blog/2023/02/orchestration-vs-choreography/>

[^9]: Hawkin, T. (2022). *Microservice Orchestration vs Choreography: How Event-Driven Architecture Helps Decouple Your App*. <https://dev.to/thawkin3/microservice-orchestration-vs-choreography-how-event-driven-architecture-helps-decouple-your-app-4a6b>

See: [Specification](#) → [Automation Mesh](#) | [Automation Mesh Diagram](#)

Layer 5: Organizational Layer (Motor Cortex)

Biological Analogy: Motor cortex coordinating complex movements

Function: Human team structures optimized for AI-native collaboration

Organizational Patterns:

1. Squads

- **Purpose:** Cross-functional product/feature teams
- **Size:** 5-9 people (Dunbar's limit for tight collaboration)
- **Structure:** Product Manager, Engineers, Designer, Data Analyst
- **AI Integration:** Embedded agents for specific squad functions
- **Ownership:** Business service accountability (P&L responsibility)
- **Lifecycle:** Persistent teams aligned to long-term product areas

2. Pools

- **Purpose:** Flexible specialist communities supporting multiple squads
- **Examples:** Data Science Pool, Security Pool, UX Research Pool
- **Model:** Pull-based engagement (squads request support)
- **AI Integration:** Pool-specific specialized agents
- **Governance:** Community lead coordinates allocation

3. Operations

- **Purpose:** Stable, repeatable business processes
- **Examples:** Payroll, Compliance, Customer Support
- **Model:** High automation (80%+ AI-driven)
- **Human Role:** Exception handling, oversight, continuous improvement
- **Metrics:** Throughput, error rate, cycle time

See: [Specification → Organizational Layer](#)

Layer 6: Governance & Ethics (Prefrontal Cortex)

Biological Analogy: Prefrontal cortex providing judgment and ethical reasoning

Function: Decision oversight ensuring alignment with values and compliance

Components:

- RFC (Request for Comments) process for major decisions
- ADR (Architecture Decision Records) documenting choices
- Ethical review board (human + AI advisors)
- Compliance monitoring (SOC2, GDPR, HIPAA, etc.)
- Incident response protocols

Key Mechanisms:

- **Impact Analysis:** Assess risks before changes

- **Approval Workflows:** Tiered authorization based on risk
- **Audit Trails:** Complete decision lineage
- **Feedback Loops:** Retrospectives driving improvement

Governance Research Validation:

SOLID.AI's governance approach aligns with emerging AI governance frameworks. Eisenberg et al. (2023)[^10] demonstrate systematic approaches to AI oversight across industries. Deloitte research (2024)[^11] highlights the critical need for transparent, auditable AI systems with human oversight for high-stakes decisions. The Governance Institute (2024)[^12] emphasizes that effective AI governance requires both automated compliance monitoring and human judgment for ethical boundaries—exactly the hybrid model SOLID.AI implements through Layer 6.

[^10]: Eisenberg, J. S., Pauwels, E., Guan, J., & Li, B. (2023). *Evaluation & Monitoring: A Research Blueprint for AI Risk Management in Practice*. arXiv. <https://arxiv.org/abs/2308.08700>

[^11]: Deloitte & Deeploy (2024). *Implementing AI Governance: A Practical Guide*. <https://www2.deloitte.com/content/dam/Deloitte/nl/Documents/deloitte-analytics/deloitte-nl-ai-deploy-report-ai-governance.pdf>

[^12]: Governance Institute of Australia (2024). *AI Oversight: What Directors Need to Know*. <https://www.governanceinstitute.com.au/resources/news/2024/ai-oversight-what-directors-need-to-know/>

See: [Governance → Implementation](#)

Organizational Scalability Model

Implementation Note: The scalability projections below are based on the Midora business plan thesis, where SOLID.AI is being applied from founding to validate this organizational model. These are strategic projections, not measured results. Midora is building the company from absolute zero using this framework, and actual performance data will be published as the implementation matures.

SOLID.AI targets exponential growth through AI multiplication:

Traditional Organization (Reference Model):

Revenue: \$10M → \$50M (+400%)

Headcount: 100 → 500 people (+400%)

Ratio: 1:1 scaling

AI-Native Organization (Projected SOLID.AI Model):
Revenue: \$10M → \$50M (+400%)
Headcount: 100 → 150 people (+50%)
AI Agents: 0 → 350 equivalent roles
Ratio: 1:0.5 scaling (humans), 1:3.5 (AI multiplication)

Projected Economic Case:

- **Traditional \$50M Company:** $500 \text{ employees} \times \$100K = \$50M \text{ payroll}$ (100% of revenue)
- **AI-Native \$50M Company (Target):** $150 \text{ employees} \times \$100K = \$15M \text{ payroll}$ (30% of revenue)
- **Projected Savings:** \$35M/year reallocated to R&D, market expansion, or profit
- **Quality Targets:** Error rates <1% (vs. 5-10% traditional), faster time-to-market

Scalability Comparison Table

Metric	Traditional Org	SOLID.AI (Projected)	Difference
Revenue Growth	\$10M → \$50M (+400%)	\$10M → \$50M (+400%)	Same growth target
Headcount Growth	100 → 500 people (+400%)	100 → 150 people (+50%)	-70% headcount
AI Agent Roles	0 agents	350 equivalent roles	+350 AI roles
Payroll Cost	\$50M (100% of revenue)	\$15M (30% of revenue)	-\$35M savings
Cost Efficiency	1:1 revenue-to-payroll	3.3:1 revenue-to-payroll	3.3x improvement
Error Rate	5-10% (manual processes)	<1% (automated quality)	5-10x improvement
Time-to-Market	Months (waterfall cycles)	Weeks (AI-accelerated)	4-10x faster
Scaling Ratio	Linear (1:1)	Exponential (1:3.5 AI multiplication)	Sublinear scaling

Note: These projections represent the Midora business plan thesis targets. Actual metrics will be published as the implementation matures in production.

Research Evidence: McKinsey Global Institute projects \$2.9 trillion in value creation through redesigning work around human-AI skill partnerships, not isolated task automation[^1]. EY

research explicitly validates "decoupling growth from headcount" and "non-linear productivity" through systematic AI integration[^2]—providing economic foundation for SOLID.AI's scalability model. McKinsey further estimates \$4.4 trillion in productivity gains when work is redesigned around "superagency" (humans supported by AI agents and automation)[^3].

[^1]: McKinsey Global Institute (2025). *Agents, Robots, and Us: Skill Partnerships in the Age of AI.* <https://www.mckinsey.com/mgi/our-research/agents-robots-and-us-skill-partnerships-in-the-age-of-ai>

[^2]: EY (2024). *AI: Ideation to Impact White Paper.* https://www.ey.com/content/dam/ey-unified-site/ey-co_m/en-in/insights/ai/documents/ey-ai-deation-to-impact.pdf

[^3]: McKinsey & Company (2025). *Superagency in the Workplace.* <https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/superagency-in-the-workplace-empowering-people-to-unlock-ais-full-potential-at-work>

Navigation: [← Abstract](#) | [Specification →](#) | [Diagrams](#)

Specification

Status: Version: 1.0

1. Core Entities

1.1 Actor

Definition: A human participant with decision-making authority and accountability within the system.

Attributes:

- `actor_id`: Unique identifier (UUID)
- `role`: Organizational role (e.g., Product Manager, Compliance Officer)
- `authority_level`: Decision boundary scope (tactical, strategic, governance)
- `authentication_context`: Identity verification state
- `session_metadata`: Active context and preferences

Constraints:

- MUST have unique identity across all system boundaries
- MUST be traceable through audit logs
- MUST operate within defined authority boundaries
- MAY delegate execution to AI Agents but CANNOT delegate accountability

Example:

```
actor:  
  actor_id: "a7f3c8b1-4e5d-6f7a-8b9c-0d1e2f3a4b5c"  
  role: "Product Manager"  
  authority_level: "strategic"  
  authentication_context:  
    method: "SSO"  
    verified_at: "2025-11-29T14:30:00Z"  
  session_metadata:  
    workspace: "Q4-Planning"  
    active_context: ["sales-analysis", "budget-review"]
```

1.3 Event

Definition: A state change or occurrence within the system that triggers downstream processing.

Attributes:

- event_id: Unique identifier (UUID)
- event_type: Classification (business, system, governance, audit)
- timestamp: ISO 8601 timestamp with timezone
- source: Originating entity (Actor, AI Agent, External System)
- payload: Event data conforming to schema
- correlation_id: Parent event or transaction identifier
- causation_chain: Full lineage of triggering events

Constraints:

- MUST be immutable after creation
- MUST include complete causation chain
- MUST be persisted to event store
- MUST propagate through Automation Mesh
- MAY trigger zero or more downstream Actions

Example:

```
event:  
  event_id: "evt-2025-11-29-14-30-001"  
  event_type: "business"  
  timestamp: "2025-11-29T14:30:15.234Z"  
  source:  
    type: "external_system"  
    system_id: "salesforce-prod"  
  payload:  
    event_name: "opportunity_closed_won"  
    opportunity_id: "opp-2025-Q4-1234"  
    amount: 250000  
    customer_id: "cust-enterprise-456"  
  correlation_id: "txn-2025-11-29-001"  
  causation_chain:  
    - "evt-2025-11-29-14-25-001" # opportunity_updated  
    - "evt-2025-11-29-14-28-003" # approval_granted
```

1.5 Policy

Definition: A declarative rule that governs system behavior, access control, and decision-making.

Attributes:

- `policy_id`: Unique identifier (UUID)
- `policy_name`: Human-readable name
- `policy_type`: Classification (*accesscontrol*, *approvalworkflow*, *data_governance*, *compliance*)
- `scope`: Applicability (global, domain-specific, agent-specific)
- `conditions`: Logical expressions for policy activation
- `enforcement_action`: Required behavior when policy triggers
- `priority`: Execution order when multiple policies apply

Constraints:

- MUST be versioned
- MUST be auditable
- MUST support conflict resolution via priority
- MAY be overridden by governance layer
- MUST be evaluated before action execution

Example:

```
policy:  
  policy_id: "pol-budget-approval-001"  
  policy_name: "Budget Allocation Approval Workflow"  
  policy_type: "approval_workflow"  
  scope:  
    domain: "finance"  
    applies_to: [ "budget-allocation", "cost-center-transfer" ]  
  conditions:  
    - "action.amount > 50000"  
    - "action.executor.type == 'ai_agent'"  
  enforcement_action:  
    type: "require_human_approval"  
    approver_roles: [ "CFO", "Finance Director" ]  
    timeout: "4h"  
  priority: 100
```

1.7 Data Domain

Definition: A logical grouping of related data entities with consistent governance, ownership, and quality standards.

Attributes:

- domain_id: Unique identifier (UUID)
- domain_name: Human-readable name
- owner: Accountable Actor or team
- schema_registry: Data structure definitions
- quality_requirements: Validation rules and SLAs
- access_control: Authorization policies
- lineage_tracking: Data provenance metadata

Constraints:

- MUST have designated owner
- MUST define schema contracts
- MUST enforce quality requirements
- MUST maintain lineage metadata
- MAY federate across multiple storage systems

Example:

```
data_domain:  
  domain_id: "dom-sales-performance-001"  
  domain_name: "Sales Performance Analytics"  
  owner:  
    actor_id: "a7f3c8b1-4e5d-6f7a-8b9c-0d1e2f3a4b5c"  
    role: "VP Sales Operations"  
  schema_registry:  
    - entity: "opportunity"  
      version: "v2.1"  
      fields: ["id", "amount", "stage", "close_date", "probability"]  
    - entity: "sales_forecast"  
      version: "v1.3"  
      fields: ["period", "amount", "confidence", "updated_at"]  
  quality_requirements:  
    completeness: ">= 0.98"  
    freshness: "<= 60s"  
    accuracy: ">= 0.95"  
  access_control:
```

```
read: ["sales_team", "executive_team", "agent-sales-analyst-*"]
write: ["salesforce-prod", "sales_automation_agents"]
```

1.9 Service Level Objectives (SLOs)

Consolidated Performance Targets

The following table defines the target Service Level Objectives for SOLID.AI framework components. These are aspirational targets for production implementations.

Component	Metric	Target	Measurement
Data Spine	Latency (P95)	< 5s	95th percentile query response time
Data Spine	Availability	$\geq 99.9\%$	System uptime over 30-day window
Data Spine	Data Freshness	< 60s	Time from source change to availability
AI Agents	Response Latency (P95)	< 5s	95th percentile from request to recommendation
AI Agents	Accuracy	$\geq 95\%$	Correct recommendations vs. validation set
AI Agents	Explainability	100%	All decisions must include reasoning
Automation Mesh	Event Processing (P95)	< 5s	Time from event emission to action initiation
Automation Mesh	Throughput	≥ 1000 events/sec	Sustained event processing capacity
Automation Mesh	Error Rate	< 1%	Failed workflows vs. total workflows
Governance	Audit Log Completeness	100%	All actions logged with full context
Governance	Override Response Time	< 100ms	Human intervention acknowledgment
Governance	Policy Violation Detection	< 1s	Time to detect and flag violations

Notes:

- **Latency vs. Freshness:** Latency measures system response time; freshness measures data currency
- **P95:** 95th percentile - 95% of requests complete within target

- **Availability:** Measured as uptime / (uptime + downtime) over rolling 30-day period
- **These are target specifications:** Actual performance depends on implementation architecture and scale

2.2 Action Orchestration

Description: The coordination of Actions across multiple systems, respecting dependencies, trust boundaries, and approval workflows.

Behavior Specification:

- **Action Planning**
- AI Agent receives Event
- Agent generates Action plan with dependencies
- Plan evaluated against Policies and trust boundaries
- **Approval Workflow**
- If action exceeds trust boundary, generate approval_request Event
- Route approval_request to appropriate Actor
- Wait for approvalgranted or approvaldenied Event (with timeout)
- **Action Execution**
- Execute Actions in dependency order
- Log execution start, progress, and completion
- Handle failures with retry and compensation logic
- **Result Propagation**
- Generate completion Event with result payload
- Update Data Spine with outcome
- Notify downstream consumers

Guarantees:

- Actions execute transactionally where possible
- Failed actions trigger compensation or rollback
- All actions traceable to originating Event

- Approval timeouts prevent indefinite blocking

Example Orchestration:

```

orchestration:
  trigger_event: "evt-opportunity-closed-won"
  planned_actions:
    - action_id: "act-001"
      type: "update_forecast"
      executor: "agent-sales-analyst-001"
      requires_approval: false
      dependencies: []

    - action_id: "act-002"
      type: "allocate_budget"
      executor: "agent-finance-automation-001"
      requires_approval: true # Amount > $50k threshold
      dependencies: ["act-001"]
      approval_workflow:
        approver_roles: ["CFO"]
        timeout: "4h"

    - action_id: "act-003"
      type: "notify_sales_team"
      executor: "agent-notification-001"
      requires_approval: false
      dependencies: ["act-001", "act-002"]

```

2.4 Context Alignment

Description: The process of ensuring AI Agents operate with current, accurate context aligned with organizational goals and real-world state.

Behavior Specification:

- **Context Acquisition**
 - AI Agent queries Data Spine for relevant data domains
 - Agent retrieves organizational objectives from Purpose Layer
 - Agent loads applicable Policies and Governance Rules
- **Context Validation**
 - Agent verifies data freshness (within SLA: < 60s)
 - Agent checks for conflicting policies
 - Agent validates against trust boundary constraints

- **Context Application**
 - Agent reasoning incorporates context into decision-making
 - Agent generates recommendations aligned with current state
 - Agent explains how context influenced output
- **Context Drift Detection**
 - System monitors for context changes (e.g., policy updates, objective shifts)
 - Out-of-date context triggers re-evaluation
 - Agent operations suspended if context invalidated

Guarantees:

- AI Agents NEVER operate with stale context
- Context freshness validated before every decision
- Context changes trigger automatic re-alignment
- Full context snapshot logged with every action

Example Context:

```

context_snapshot:
  agent_id: "agent-sales-analyst-001"
  timestamp: "2025-11-29T14:30:00Z"
  data_spine_context:
    - domain: "sales_performance"
      freshness: "12s"
      entities: ["opportunities", "forecasts", "pipeline"]
  purpose_layer_context:
    objectives:
      - "Achieve Q4 revenue target: $2.5M"
      - "Maintain sales cycle < 30 days"
    priorities: ["revenue_growth", "customer_retention"]
  policy_context:
    applicable_policies:
      - "pol-budget-approval-001"
      - "pol-forecast-accuracy-001"
  governance_context:
    active_rules:
      - "gov-gdpr-art22-001"
      - "gov-sox-404-001"

```

3. System Guarantees

3.1 Deterministic Edges

Guarantee Statement: All decision points and state transitions in the system produce consistent, predictable outcomes given identical inputs.

Specification:

- **Idempotency:** Repeating the same Action with the same parameters produces the same result
- **Reproducibility:** Given the same Event and context, AI Agents generate identical recommendations
- **Predictability:** Policy evaluation produces consistent enforcement actions
- **Testability:** All system behaviors verifiable via automated testing

Implementation Requirements:

• **Deterministic AI Models:**

- Set temperature=0 for reproducible outputs
- Use fixed random seeds in testing
- Version-lock model references

• **Immutable Events:**

- Events never modified after creation
- Event replay produces identical downstream effects

• **Stateless Processing:**

- Actions depend only on inputs and context
- No hidden state or side effects

Verification:

```
# Example deterministic test
def test_forecast_update_deterministic():
    event = Event(type="opportunity_closed_won", payload={"amount": 250000})
    context = Context(q4_target=2500000, current_forecast=2200000)

    agent = SalesAnalystAgent(temperature=0, model="gpt-4o-2024-11")

    result1 = agent.process(event, context)
    result2 = agent.process(event, context)

    assert result1 == result2 # Deterministic output
    assert result1.new_forecast == 2450000
```

3.3 Compliance Invariants

Guarantee Statement: The system maintains continuous compliance with all applicable regulations, standards, and organizational policies under all operating conditions.

Specification:

- **Policy Enforcement:** All Policies evaluated before Action execution
- **Governance Supremacy:** Governance Rules override conflicting behaviors
- **Boundary Integrity:** No operations cross Boundaries without authorization
- **Audit Completeness:** All compliance-relevant activities logged

Implementation Requirements:

- **Policy Engine:**
 - Centralized policy evaluation before every action
 - Policy conflicts resolved via priority
 - Policy violations block execution
- **Governance Layer:**
 - Continuous monitoring of active Governance Rules
 - Real-time violation detection
 - Automatic remediation or escalation
- **Compliance Validation:**
 - Automated compliance checks (e.g., GDPR, SOX, HIPAA)
 - Regular compliance audits via external tooling
 - Compliance dashboard for real-time visibility

Supported Regulations:

- **GDPR:** Right to explanation, consent management, data minimization
- **SOX:** Financial controls, audit trails, segregation of duties
- **HIPAA:** PHI access controls, encryption, breach notification

- **ISO 27001:** Information security management
- **FedRAMP:** Cloud security for government data

Example Invariant Check:

```

compliance_check:
  rule_id: "gov-gdpr-art22-001"
  check_time: "2025-11-29T14:30:15Z"
  action_under_review:
    action_id: "act-budget-allocation-789"
    affects_individual_rights: true
    executor: "agent-finance-automation-001"
  validation_result:
    compliant: false
    violations:
      - "Action affects individual rights but lacks explainability"
      - "No human review requested (required for high-impact decisions)"
  enforcement:
    action_blocked: true
    remediation: "Require human approval from CFO"
    notification_sent_to: ["DPO", "compliance_team"]

```

4. Conformance Testing

Implementations claiming SOLID.AI compliance MUST pass the following conformance test suite:

4.1 Entity Conformance

- █ All core entities implement required attributes
- █ Entity constraints enforced at runtime
- █ Entity serialization follows specification

4.2 Behavior Conformance

- █ Event propagation maintains causation chains
- █ Action orchestration respects trust boundaries
- █ Human override latency < 100ms
- █ Context alignment validates freshness
- █ Audit trail achieves 100% coverage

4.3 Guarantee Conformance

- Deterministic edges verified via automated tests
- Traceability validated end-to-end
- Compliance invariants continuously monitored
- Observability metrics published and alertable

Conformance Certification:

Implementations passing all conformance tests receive **SOLID.AI v1 Certified** designation.

6. Version History

| Version | Date | Changes |

|-----|-----|-----|

| 1.0 | 2025-11-29 | Initial stable release |

Navigation:

- ← [Back to Whitepaper Index](#)
- [View Architecture →](#)
- [View Technical Specification →](#)

Implementation Guide

Status: Version: 1.0

This section provides detailed technical specifications for each layer of the SOLID.AI architecture.

Cognitive Layer (Layer 3)

Overview

As shown in Figure 4 (see [Diagrams](#)), AI agents operate as organizational members with defined roles, capabilities, and accountability through a continuous collaboration loop with humans.

See: [Human-AI Collaboration Loop Diagram](#) for complete interaction flow

Agent Definition Schema

```
agent:  
  id: sales_analyst_001  
  name: Sales Performance Analyzer  
  type: analytical  
  version: 2.1.0  
  
  purpose: |  
    Analyze sales pipeline data, identify trends, and provide  
    actionable recommendations to sales leadership.  
  
  capabilities:  
    - pipeline_forecasting  
    - deal_risk_assessment  
    - win_loss_analysis  
    - competitor_intelligence  
  
  data_access:  
    read:  
      - customers (all)  
      - opportunities (all)  
      - contracts (all)  
      - interactions (type="sales_call")  
    write:
```

```

    - forecasts (own)
    - recommendations (own)

interfaces:
  input:
    - slack_channel: "#sales-analytics"
    - api_endpoint: "/agents/sales_analyst"
    - scheduled_triggers: ["daily 8am", "weekly monday"]
  output:
    - slack_notifications: true
    - dashboard_updates: "sales_dashboard"
    - email_reports: sales_leadership@company.com

constraints:
  execution_time: <60 seconds
  cost_per_run: <$0.50
  accuracy_threshold: >95%

human_oversight:
  approval_required: false
  audit_frequency: weekly
  escalation_conditions:
    - forecast_deviation >20%
    - deal_risk_score >8/10

ethical_boundaries:
  - no_customer_discrimination
  - transparent_scoring_methodology
  - human_review_for_contract_termination

```

Agent Lifecycle

- **Definition:** RFC process for new agents
- **Development:** Build and test in sandbox
- **Validation:** Human review + test cases
- **Deployment:** Gradual rollout with monitoring
- **Operation:** Continuous execution + logging
- **Evolution:** Feedback-driven improvements
- **Retirement:** Deprecation with migration plan

Reasoning Patterns

Chain-of-Thought:

User Query: "Why is Q4 forecast down 15%?"

Agent Reasoning:

1. Retrieve Q4 pipeline data
2. Compare to Q3 pipeline at same point
3. Identify closed-lost deals (reasons)
4. Analyze new deal velocity (slower)
5. Assess stage progression rates (delayed)
6. Synthesize findings into explanation

Output: "Q4 forecast is down 15% due to: (1) 3 large enterprise deals slipped to Q1 (\$450K total), (2) new pipeline generation 20% below target, and (3) slower progression from Discovery → Proposal (avg 14 days vs. 9 days in Q3). Recommendation: Focus on accelerating mid-stage deals and launching Q1 demand gen campaign."

Human-AI Collaboration Model:

- AI performs analysis (speed, scale)
- Human validates conclusions (judgment)
- AI implements decisions (execution)
- Human monitors outcomes (oversight)

Organizational Layer (Layer 5)

Squad Specification

Charter Template:

```
squad:  
  name: Checkout Experience Squad  
  mission: Optimize conversion and revenue at checkout  
  
  business_service:  
    name: E-Commerce Checkout  
    metrics:  
      - conversion_rate (current: 68%, target: 75%)  
      - cart_abandonment (current: 32%, target: 25%)  
      - revenue_per_session (current: $45, target: $55)  
  
  team:  
    product_manager: alice_johnson  
    tech_lead: bob_chen  
    engineers: [carol_lopez, dave_kumar, eve_taylor]  
    designer: frank_williams  
    data_analyst: grace_martinez
```

```
ai_agents:
  - checkout_optimizer (A/B test orchestration)
  - fraud_detector (transaction risk scoring)
  - personalization_engine (offer recommendations)

dependencies:
  upstream:
    - Product Catalog Squad (inventory data)
    - Pricing Squad (promotional rules)
  downstream:
    - Order Fulfillment Squad (order handoff)
    - Customer Support Squad (checkout issues)

ceremonies:
  sprint_length: 2 weeks
  planning: Monday 9am
  daily_standup: Daily 10am (15 min)
  review: Friday 2pm
  retrospective: Friday 3pm

decision_authority:
  autonomous: [UI changes, A/B tests, bug fixes]
  requires_approval: [pricing strategy, payment provider]
  forbidden: [PCI compliance changes without Security]
```

Navigation: [← Architecture](#) | [Governance →](#) | [Diagrams](#)

Governance

Status: Version: 1.0

Implementation Methodology

SOLID.AI transformation follows a phased approach balancing speed with organizational change management. As shown in Figure 1 (see [Diagrams](#)), the complete architecture is built incrementally, starting with the Data Spine (Figure 3) and Cognitive Layer (Figure 4), then scaling the Automation Mesh (Figure 2) across the organization.

Three-Phase Roadmap

Phase 1: Foundation (Months 1-3)

Objectives:

- Establish Data Spine infrastructure
- Define Purpose Layer (mission, values, OKRs)
- Select pilot business service
- Form first AI-native squad

Deliverables:

- [] Canonical data models documented
- [] Data contracts between 3+ systems
- [] First AI agent deployed (low-risk use case)
- [] RFC/ADR governance process established
- [] Ethical review board formed

Success Metrics:

- Data Spine operational (availability $\geq 99.9\%$)

- P95 latency < 5s for data propagation
- Data freshness < 60s for real-time entities
- First agent achieving >90% accuracy
- Zero ethical violations

Pilot Candidates:

- Sales pipeline analysis (low risk, high value)
- Customer support ticket routing
- Invoice processing automation
- Marketing campaign performance analysis

Phase 2: Pilot & Learn (Months 4-9)

Objectives:

- Scale to 3-5 squads across functions
- Deploy 10-15 production AI agents
- Validate organizational patterns
- Refine governance processes

Deliverables:

- [] 3 business services AI-native
- [] Cross-functional squad coordination proven
- [] Agent marketplace established (reusable agents)
- [] Observability dashboards operational
- [] First retrospective-driven improvements

Success Metrics:

- 50% reduction in cycle time (pilot services)
- 80% automation rate for operational tasks
- Employee satisfaction >4.0/5.0
- Zero compliance incidents

Common Challenges:

- Resistance to change (address with training)
- Data quality issues (invest in cleanup)
- Integration complexity (prioritize key systems)
- Unclear roles (define RACI matrices)

Phase 3: Scale (Months 10-24)

Objectives:

- Whole-organization transformation
- 50+ AI agents in production
- All functions operating AI-native
- Self-sustaining continuous improvement

Deliverables:

- [] 100% business services AI-enabled
- [] Agent autonomy increasing (80%+ decisions)
- [] Organizational scalability demonstrated
- [] Documented playbooks for new entrants
- [] Open-source contributions to framework

Success Metrics:

- 10x improvement in time-to-market
- Revenue growth without linear headcount scaling
- <1% error rates across processes
- Industry recognition (case studies, awards)

Compliance Management

Regulatory Frameworks

SOLID.AI supports compliance with:

Framework Scope Key Requirements
----- ----- -----
GDPR EU data protection Consent, data minimization, right to erasure
CCPA California privacy Disclosure, opt-out, non-discrimination
SOC 2 Security controls Access control, encryption, audit logs
HIPAA Healthcare data PHI protection, access logging, encryption
ISO 27001 Information security Risk assessment, incident response
FedRAMP US government cloud Enhanced security controls, continuous monitoring

Compliance Architecture

Data Classification:

```
data_classification:  
  public:  
    examples: [marketing_content, blog_posts]  
    encryption: optional  
    access: all  
  
  internal:  
    examples: [roadmaps, financial_models]  
    encryption: required  
    access: employees_only  
  
  confidential:  
    examples: [customer_contracts, employee_salaries]  
    encryption: required (AES-256)  
    access: role_based  
    audit: all_access_logged  
  
  restricted:  
    examples: [PHI, PII, financial_transactions]  
    encryption: required (AES-256 + tokenization)  
    access: explicit_approval  
    audit: all_access_logged + reviewed  
    retention: auto_delete_after_90_days
```

Agent Compliance Controls:

```
agent: customer_support_assistant
compliance:
  data_access:
    - customer_name (public)
    - email (confidential, masked: j***@example.com)
    - order_history (confidential)
    - payment_info (FORBIDDEN - restricted)

  retention:
    conversation_logs: 90_days
    sensitive_data: 30_days
    audit_trail: 7_years

  encryption:
    in_transit: TLS 1.3
    at_rest: AES-256

  monitoring:
    access_logging: enabled
    anomaly_detection: enabled
    compliance_alerts: pii_exposure, unauthorized_access
```

Continuous Improvement

Feedback Loops

Agent Performance Review (Weekly):

- Accuracy metrics vs. baseline
- Cost per execution
- User satisfaction ratings
- Error analysis

Squad Retrospective (Biweekly):

- What went well?
- What needs improvement?
- Action items (captured as RFC/ADR)

Organizational Health Check (Quarterly):

- Employee engagement survey
- AI trust metrics
- Ethical incident review
- Scalability assessment

Annual Framework Audit:

- Purpose Layer relevance
- Architecture evolution needs
- Governance effectiveness
- Industry benchmark comparison

Non-Linear Productivity & Economic Impact

SOLID.AI's scalability projections are grounded in emerging research demonstrating that systematic AI integration enables organizations to decouple revenue growth from headcount expansion—fundamentally changing traditional linear economic models.

McKinsey & Company. (2023). *The Economic Potential of Generative AI: The Next Productivity Frontier.*

<https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/the-economic-potential-of-generative-ai-the-next-productivity-frontier>

Estimates generative AI could add \$2.6-4.4 trillion annually to the global economy, increasing total AI impact by 15-40%. Provides economic validation for SOLID.AI's projection of exponential productivity gains through systematic AI integration.

McKinsey Global Institute. (2025). *Agents, Robots, and Us: Skill Partnerships in the Age of AI.*

<https://www.mckinsey.com/mgi/our-research/agents-robots-and-us-skill-partnerships-in-the-age-of-ai>

Projects \$2.9 trillion in value creation through redesigning work around partnerships between humans, AI agents, and automation—not isolated task automation. Directly supports SOLID.AI's organizational scalability model showing revenue growth decoupled from headcount.

McKinsey & Company. (2025). Superagency in the Workplace: Empowering People to Unlock AI's Full Potential at Work.

<https://www.mckinsey.com/capabilities/tech-and-ai/our-insights/superagency-in-the-workplace-empowering-people-to-unlock-ais-full-potential-at-work>

Introduces the concept of "superagency"—people supported by AI agents and automation—estimating \$4.4 trillion in productivity gains when work is redesigned around human-AI collaboration.

EY. (2024). AI: Ideation to Impact White Paper.

https://www.ey.com/content/dam/ey-unified-site/ey-com/en-in/insights/ai/documents/ey-ai-deation-to-impac_t.pdf

Explicitly states AI enables "decoupling growth from headcount" and "non-linear productivity"—the exact economic model underlying SOLID.AI's scalability projections (3.3:1 revenue-to-payroll ratio vs. traditional 1:1).

People Managing People. (2024). AI Case Studies in Operations and Business Process Outsourcing.

<https://peoplemanagingpeople.com/hr-strategy/examples-of-ai-in-hr/>

Provides real-world examples of companies achieving non-linear scaling: maintaining or growing workload without proportional headcount increases through systematic AI integration.

Navigation: [← Specification](#) | [← Abstract](#) | [☰ Diagrams](#)

References & Further Reading

This section provides academic and industry research supporting SOLID.AI's architectural decisions, economic projections, and organizational patterns.

AI-Native Organizations

Wile, R., & Wilson, H. J. (2019). Building the AI-Powered Organization. Harvard Business Review.

<https://hbr.org/2019/07/building-the-ai-powered-organization>

Demonstrates that organizational structure, data infrastructure, and processes—not technology—are the primary bottlenecks to AI adoption. This directly validates SOLID.AI's focus on Data Spine, Automation Mesh, and Governance as foundational layers.

Harvard Business School Online. (2025). How to Architect an AI-Native Business.

<https://online.hbs.edu/blog/post/ai-native>

Examines companies designed from inception as AI-native, with AI embedded in strategic decisions and operational processes. Aligns with SOLID.AI's concept of the "natively cognitive organization."

Interloom. (2024). AI-Native Organizations.

<https://www.interloom.com/en/blog/ai-native-organizations>

Defines AI-native organizations as those that capture tacit knowledge, embed agents directly into workflows, and enable real-time coordination—an operational description of SOLID.AI's Automation Mesh and Cognitive Layer integration.

Ema. (2024). Understanding the Concept of AI Native and its Impact on Business.

<https://www.ema.co/additional-blogs/addition-blogs/understanding-the-concept-of-ai-native-and-its-impact-on-business>

Defines AI-native as having AI at the center of architecture, decisions, and culture—consistent with SOLID.AI's Purpose Layer and Cognitive Layer design.

Event-Driven Automation & Orchestration

Camunda. (2023). Orchestration vs. Choreography in Microservices.

<https://camunda.com/blog/2023/02/orchestration-vs-choreography/>

Explains advantages and tradeoffs of centralized orchestration vs. event-based choreography, and strategies for combining both approaches—precisely what SOLID.AI implements with

Automation Mesh as event mesh + orchestration fabric.

Hawkin, T. (2022). *Microservice Orchestration vs Choreography: How Event-Driven Architecture Helps Decouple Your App.* DEV Community.

<https://dev.to/thawkin3/microservice-orchestration-vs-choreography-how-event-driven-architecture-helps-de-couple-your-app-4a6b>

Demonstrates how event-driven architectures decouple services, provide resilience, and enable scalability—the same principles SOLID.AI applies in coupling AI agents, business services, and human workflows through the Automation Mesh.

AI Governance, Risk & Compliance

Eisenberg, D., et al. (2025). *The Unified Control Framework: Establishing a Common Foundation for Enterprise AI Governance, Risk Management and Regulatory Compliance.* arXiv:2503.05937.

<https://arxiv.org/abs/2503.05937>

Proposes a unified framework integrating AI governance, risk management, and compliance into enterprise architecture—validates SOLID.AI's Governance Layer approach with embedded controls, audit trails, and ethical review processes.

Deepoly, Deloitte, et al. (2025). *AI Governance & Control Framework White Paper.*

<https://deepoly.ml/white-paper-ai-governance-control-framework/>

Defines practical roadmap for implementing governance throughout the AI lifecycle without blocking innovation—reinforces SOLID.AI's integration of governance into Automation Mesh, Data Spine, and agent deployment pipelines.

Governance Institute of Australia. (2024). *White Paper on AI Governance.*

<https://www.governanceinstitute.com.au/app/uploads/2024/09/GovInst-AI-Whitepaper.pdf>

Emphasizes accountability, transparency, and risk management as fundamental for safe AI adoption at scale—all explicitly addressed in SOLID.AI's RFC/ADR processes, ethical review boards, and compliance monitoring.

How to Cite SOLID.AI

If you use SOLID.AI in your research or project, please cite:

```
@dataset{solidai_zenodo_2025,
  title      = {SOLID.AI Framework – Whitepaper v1.0},
  author     = {Freitas, Gustavo},
  year       = 2025,
  month      = november,
  publisher  = {Zenodo},
  doi        = 10.5281/zenodo.17765515,
  url        = https://zenodo.org/records/17765515
}
```

Citation Format

APA:

Freitas, G. (2025). SOLID.AI Framework – Whitepaper v1.0 [Dataset]. Zenodo. <https://doi.org/10.5281/zenodo.17765515>

IEEE:

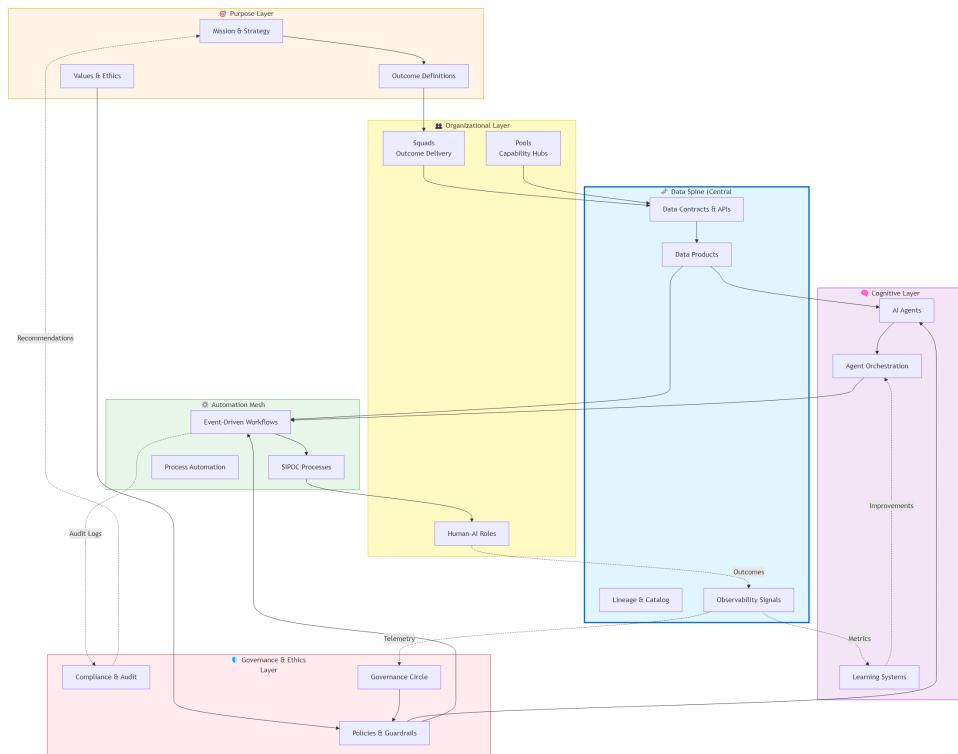
G. Freitas, "SOLID.AI Framework – Whitepaper v1.0," Zenodo, Nov. 2025.
doi: 10.5281/zenodo.17765515

Navigation: [← Diagrams](#) | [Abstract →](#)

Whitepaper Diagrams

Status: Version: 1.0

1. SOLID.AI Architecture Layer Model



<figure markdown>

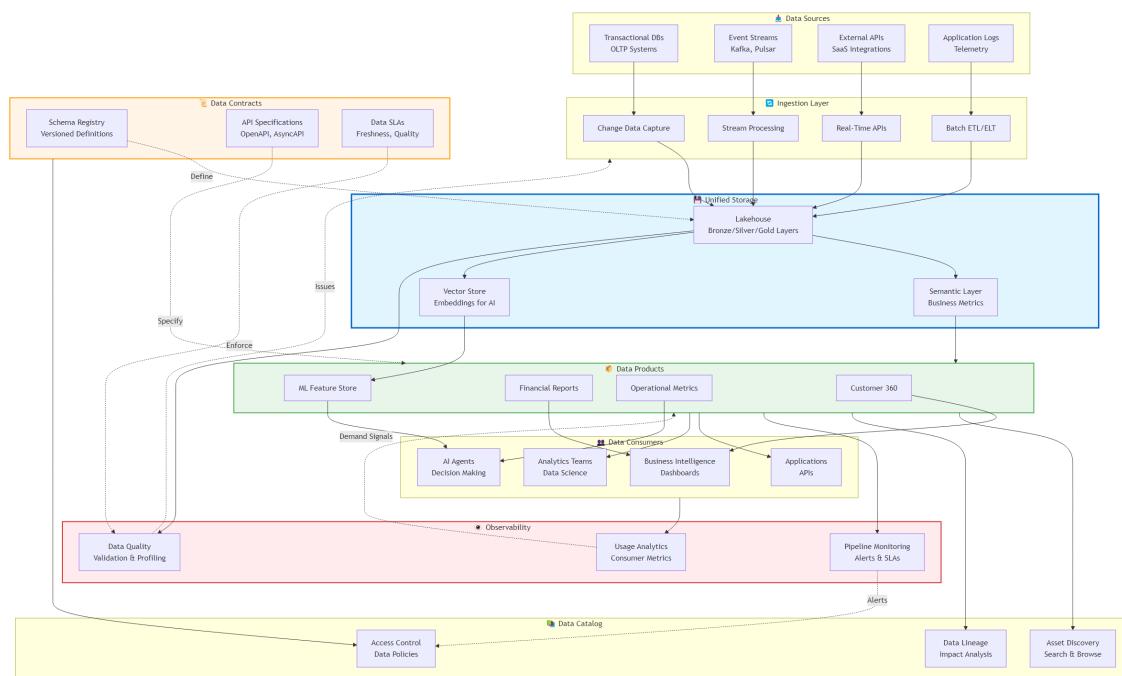
<figcaption>Figure 1 — SOLID.AI Six-Layer Architecture

Overview of the structural layers from alignment to governance, establishing the foundation for hybrid intelligent organizations. The six-layer architecture creates an organizational nervous system where the Purpose Layer (DNA) defines immutable identity, the Data Spine (sensory nerves) provides real-time information, the Cognitive Layer (brain) generates insights, the Automation Mesh (motor neurons) executes processes, the Organizational Layer (motor cortex) coordinates human teams, and the Governance Layer (prefrontal cortex) ensures ethical oversight.</figcaption>

</figure>

Reference Implementation: See [Midora Topology \(ADR-0003\)](#) for a concrete mapping of this diagram into a real AI-native education platform.

3. SOLID.AI Data Spine Topology



<figure markdown>

<figcaption>Figure 3 — Data Spine Domain Model

Unified data backbone enabling clean, derived, and real-time data flows across the organization. The Data Spine ingests data from multiple sources (CRM, ERP, Support, Product, HR) via CDC/APIs/webhooks, streams events through Kafka, maps to canonical entity models, validates quality, stores in PostgreSQL (transactional) and Data Warehouse (analytics), and serves all consumers with < 5 second latency and 99.9% uptime SLA. **Target SLO:** P95 latency < 5 s, availability $\geq 99.9\%$, data freshness < 60 s for real-time entities.</figcaption>

</figure>

Reference Implementation: See [Midora Topology \(ADR-0003\)](#) for a concrete mapping of this diagram into a real AI-native education platform.

Diagram Usage Guidelines

In Academic Citations

When referencing these diagrams in papers:

"Figure 1 shows the SOLID.AI six-layer architecture (Freitas, 2025), where each layer serves a distinct biological function in the organizational nervous system."

In Implementation

These diagrams should be used during:

- **Executive presentations** - Use Layer Model to explain transformation scope
- **Technical architecture reviews** - Reference Data Spine and Automation Mesh for infrastructure design
- **Team onboarding** - Show Human-AI Collaboration Loop to clarify roles
- **Vendor evaluations** - Map vendor capabilities to specific layers

Diagram Formats

All diagrams are available in multiple formats:

- **Mermaid (source)** - Editable, version-controlled .mmd files
- **SVG (web)** - Rendered automatically in browser, scalable for presentations
- **PNG (print)** - High-resolution exports for documentation and papers
- **PDF (publication)** - Vector format for academic submissions

Navigation: [← Governance](#) | [Abstract →](#)